

Day 4 programs

1. Find the performance of mojo programming with optimization and vectorization for matrix transpose problem statement. [Note: size of matrix: 128 x 128]

```
MOJO > DAY-4 > 1.mojo
1  from sys.intrinsics import sizeof
2  from memory import memset_zero
3  from time import now
4
5  struct Matrix:
6      var data: DTypePointer[DType.float32]
7      var rows: Int
8      var cols: Int
9
10     fn __init__(inout self, rows: Int, cols: Int):
11         self.rows = rows
12         self.cols = cols
13         self.data = DTypePointer[DType.float32].alloc(rows * cols)
14         memset_zero(self.data, rows * cols)
15
16     fn __del__(owned self):
17         self.data.free()
18
19     fn __getitem__(self, row: Int, col: Int) -> Float32:
20         return self.data.load(row * self.cols + col)
21
22     fn __setitem__(self, row: Int, col: Int, val: Float32):
23         self.data.store(row * self.cols + col, val)
24
25     fn min(a: Int, b: Int) -> Int:
26         if a < b:
27             return a
28         return b
```

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS

```
chocku@imperio:/mnt/c/Users/chock/OneDrive/Desktop/MOJO$ cd DAY-4
chocku@imperio:/mnt/c/Users/chock/OneDrive/Desktop/MOJO/DAY-4$ mojo 1.mojo
Matrix size: 128 x 128
Unoptimized transpose time: 2.999999999999997e-08 seconds
Optimized transpose time: 2e-08 seconds
Speedup: 1.4999999999999998 x
chocku@imperio:/mnt/c/Users/chock/OneDrive/Desktop/MOJO/DAY-4$
```

2. Find the performance of mojo programming with parallelization for matrix addition and subtraction in a single program. [Note: size of matrix: 128 x 128]

```
In [2]: from benchmark import Unit
        from sys.intrinsics import strided_load
        from math import CeilDivisibleRaising
        from memory import memset_zero
        from memory.unsafe import DTypePointer
        from random import rand, random_float64
        from sys.info import simdwidthof
        from runtime.llcl import Runtime
        from algorithm import parallelize
        import time
```

```
In [3]: alias type = DType.float32

        struct Matrix[rows: Int, cols: Int]:
            var data: DTypePointer[type]

            # Initialize zeroing all values
            fn __init__(inout self):
                self.data = DTypePointer[type].alloc(rows * cols)
                memset_zero(self.data, rows * cols)

            # Initialize taking a pointer, don't set any elements
            fn __init__(inout self, data: DTypePointer[type]):
                self.data = data

            # Initialize with random values
            @staticmethod
            fn rand() -> Self:
                var data = DTypePointer[type].alloc(rows * cols)
                rand(data, rows * cols)
                return Self(data)

            fn __getitem__(self, y: Int, x: Int) -> Scalar[type]:
                return self.load[1](y, x)

            fn __setitem__(self, y: Int, x: Int, val: Scalar[type]):
                self.store[1](y, x, val)

            fn load[nelts: Int](self, y: Int, x: Int) -> SIMD[type, nelts]:
                return self.data.load[width=nelts](y * self.cols + x)

            fn store[nelts: Int](self, y: Int, x: Int, val: SIMD[type, nelts]):
                return self.data.store[width=nelts](y * self.cols + x, val)
```

```
In [4]: alias nelts = simdwidthof[DType.float32]() * 2
```

```
In [5]: from algorithm import vectorize
        from algorithm import parallelize

        #write similar code for fn matadd_parallelized
        fn matadd_parallelized(C: Matrix, A: Matrix, B: Matrix):
            @parameter
            fn calc_row(m: Int):
                for n in range(A.cols):
                    @parameter
                    fn add[nelts : Int](n : Int):
                        C.store[nelts](m,n, A[m,n] + B[m,n])
                        vectorize[add, nelts, size = C.cols]()
                    parallelize[calc_row](C.rows, C.rows)
            fn matsub_parallelized(C: Matrix, A: Matrix, B: Matrix):
                @parameter
                fn calc_row(m: Int):
                    for n in range(A.cols):
                        @parameter
                        fn sub[nelts : Int](n : Int):
                            C.store[nelts](m,n, A[m,n] - B[m,n])
                            vectorize[sub, nelts, size = C.cols]()
                        parallelize[calc_row](C.rows, C.rows)
```

```
In [6]: alias N = 128
```

3. WAP using multi-layer perceptron model to implement the functionality of XOR gate by using 3 different activation functions and show the change in error.

MOJO > DAY-4 > 3.mojo

```
1  from math import exp, tanh
2  from random import seed, random_float64
3  from time import now
4
5
6  struct Array:
7      var data: Pointer[Float64]
8      var size: Int
9
10     fn __init__(inout self, size: Int):
11         self.size = size
12         self.data = Pointer[Float64].alloc(self.size)
13
14     fn __init__(inout self, size: Int, default_value: Float64):
15         self.size = size
16         self.data = Pointer[Float64].alloc(self.size)
17         for i in range(self.size):
18             self.data.store(i, default_value)
19
20     fn __copyinit__(inout self, copy: Array):
21         self.size = copy.size
22         self.data = Pointer[Float64].alloc(self.size)
23         for i in range(self.size):
24             self.data.store(i, copy[i])
25
26     fn __getitem__(self, i: Int) -> Float64:
27         return self.data.load(i)
28
```

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS JUPYTER

```
Epoch 9800 loss = 0.36784116604498074
Epoch 9900 loss = 0.36784116604498074
```

Predictions:

```
0 0 1
0 1 0
1 0 1
1 1 0
```

chocku@imperio: /mnt/c/Users/chock/OneDrive/Desktop/MOJO/DAY-4\$

